

Aloha-HE

A Low-Area Hardware Accelerator for
Client-Side Operations in Homomorphic Encryption

Florian Krieger, Florian Hirner, Ahmet Can Mert, Sujoy Sinha Roy
Institute of Applied Information Processing and Communications, TU Graz

January 31, 2024

Outline

- 1 Motivation for Homomorphic Encryption
- 2 Our Proposed Design
- 3 Implementation Results & Comparison
- 4 Conclusion

Table of Contents

1 Motivation for Homomorphic Encryption

2 Our Proposed Design

3 Implementation Results & Comparison

4 Conclusion

Homomorphic Encryption (in short HE)

- Special encryption technique
- Computations directly on encrypted data
 - ➔ Confidential cloud computing
 - ➔ Sensitive data processing on untrusted third parties
 - ➔ Novel opportunities


Homomorphic Encryption (in short HE)

- Special encryption technique
- Computations directly on encrypted data
 - ➔ Confidential cloud computing
 - ➔ Sensitive data processing on untrusted third parties
 - ➔ Novel opportunities

Homomorphic Encryption (in short HE)


- Special encryption technique
- Computations directly on encrypted data
 - ➔ Confidential cloud computing
 - ➔ Sensitive data processing on untrusted third parties
 - ➔ Novel opportunities

Homomorphic Encryption: An Example




- Encrypts sensitive data
- Sends it to cloud
- Computes on encrypted data
 - E.g.: AI application
- Returns result to client

Homomorphic Encryption: An Example




- Encrypts sensitive data
- Sends it to cloud
- Computes on encrypted data
 - E.g.: AI application
- Returns result to client

Homomorphic Encryption: An Example




- Encrypts sensitive data
- Sends it to cloud
- Computes on encrypted data
 - E.g.: AI application
- Returns result to client

Homomorphic Encryption: An Example



- Encrypts sensitive data
- Sends it to cloud
- Computes on encrypted data
 - E.g.: AI application
- Returns result to client

Homomorphic Encryption: An Example



- Encrypts sensitive data
- Sends it to cloud
 - Computes on encrypted data
 - E.g.: AI application
- Returns result to client

Homomorphic Encryption: The Drawback

- Huge computational overhead
 - ➔ Large polynomials
- Simple computations take seconds
 - ➔ Hardens practical adoption

Homomorphic Encryption: The Drawback

- Huge computational overhead
 - ➔ Large polynomials
- Simple computations take seconds
 - ➔ Hardens practical adoption

Homomorphic Encryption: The Drawback

- Huge computational overhead
 - ➔ Large polynomials
- Simple computations take seconds
 - ➔ Hardens practical adoption

Homomorphic Encryption: The Drawback

- Huge computational overhead
 - ➔ Large polynomials
- Simple computations take seconds
 - ➔ Hardens practical adoption

Homomorphic Encryption: The Drawback


- Huge computational overhead
 - ➔ Large polynomials
- Simple computations take seconds
 - ➔ Hardens practical adoption



Hardware acceleration!


State of the art in accelerating HE

State of the art in accelerating HE




- Strong focus on server-side operations
 - ➡ Usually the limiting factor
- **But:** Client-side is also crucial in constrained scenarios

State of the art in accelerating HE




- Strong focus on server-side operations
 - ➔ Usually the limiting factor
- **But:** Client-side is also crucial in constrained scenarios

State of the art in accelerating HE




- Strong focus on server-side operations
 - ➔ Usually the limiting factor
- **But:** Client-side is also crucial in constrained scenarios

State of the art in accelerating HE




- Consider IoT application with HE:
 - Timely tasks require fast encryption/decryption
 - Limited energy budget
 - Low-cost devices
- Only few works targeting client-side operations
 - ➔ Insufficient hardware support

State of the art in accelerating HE




- Consider IoT application with HE:
 - Timely tasks require fast encryption/decryption
 - Limited energy budget
 - Low-cost devices
- Only few works targeting client-side operations
 - ➔ Insufficient hardware support

State of the art in accelerating HE




- Consider IoT application with HE:
 - Timely tasks require fast encryption/decryption
 - Limited energy budget
 - Low-cost devices
- Only few works targeting client-side operations
 - ➔ Insufficient hardware support

State of the art in accelerating HE




- Consider IoT application with HE:
 - Timely tasks require fast encryption/decryption
 - Limited energy budget
 - Low-cost devices
- Only few works targeting client-side operations
 - ➔ Insufficient hardware support

State of the art in accelerating HE



- Consider IoT application with HE:
 - Timely tasks require fast encryption/decryption
 - Limited energy budget
 - Low-cost devices
- Only few works targeting client-side operations
 - ➔ Insufficient hardware support

State of the art in accelerating HE



- Consider IoT application with HE:
 - Timely tasks require **fast encryption/decryption**
 - Limited **energy budget**
 - **Low-cost devices**
- Only few works targeting client-side operations
 - ➔ Insufficient hardware support

} **Motivation for Aloha-HE**

The CKKS Scheme

The CKKS Scheme

- State-of-the-art homomorphic encryption scheme
- Allows computations on complex numbers (\mathbb{C})
 - ➔ Required in machine learning applications
- Software support such as in Microsoft SEAL library

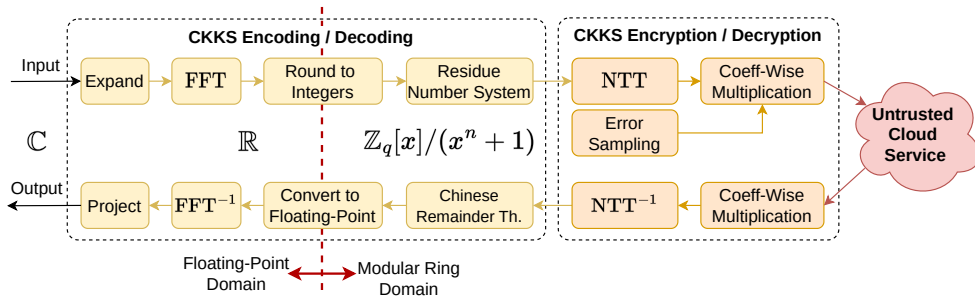
The CKKS Scheme

- State-of-the-art homomorphic encryption scheme
- Allows computations on complex numbers (\mathbb{C})
 - ➔ Required in machine learning applications
- Software support such as in Microsoft SEAL library

The CKKS Scheme

- State-of-the-art homomorphic encryption scheme
- Allows computations on complex numbers (\mathbb{C})
 - ➔ Required in machine learning applications
- Software support such as in Microsoft SEAL library

The CKKS Scheme: Operations involved



The CKKS Scheme: Operations involved

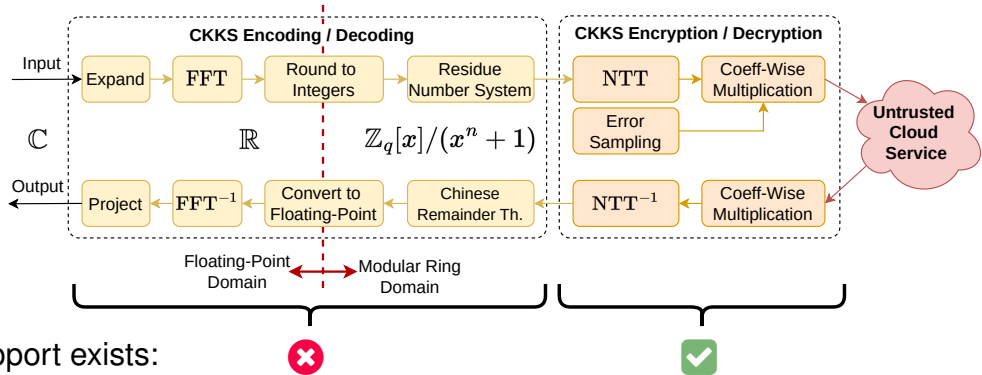


Table of Contents


1 Motivation for Homomorphic Encryption

2 Our Proposed Design

3 Implementation Results & Comparison


4 Conclusion

Overall Architecture




- CKKS coprocessor performing all client-side operations
- DMA for data streaming
- CPU controls DMA and coprocessor
- Instruction set based

Overall Architecture



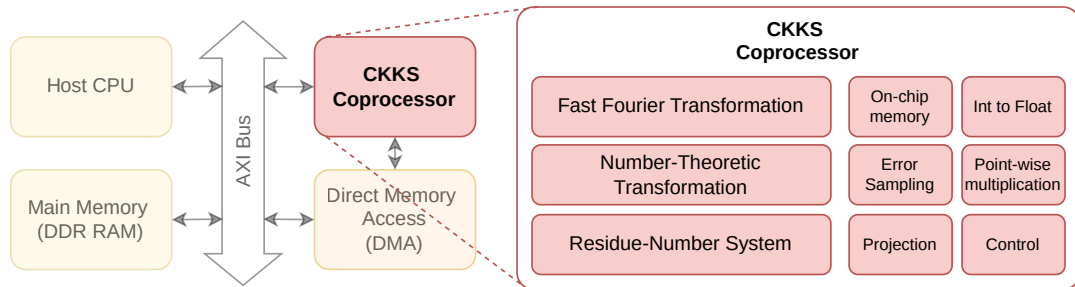
- CKKS coprocessor performing all client-side operations
- DMA for data streaming
- CPU controls DMA and coprocessor
 - Instruction set based

Overall Architecture

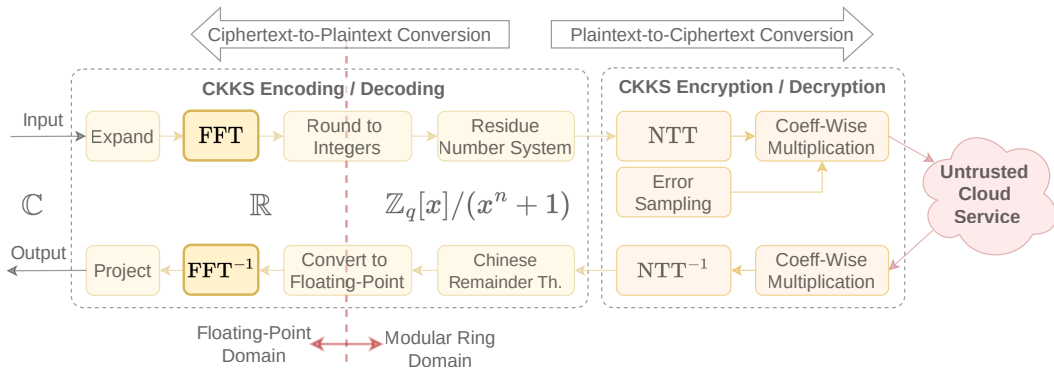


- CKKS coprocessor performing all client-side operations
- DMA for data streaming
- CPU controls DMA and coprocessor
 - Instruction set based

Overall Architecture




Fast Fourier Transformation




Fast Fourier Transformation: Data Flow

- Basic operation: Butterfly
- Iterated within $\log_2(n)$ stages
- Requires twiddle factors ω^i
 - Roots of unity in \mathbb{C}




Fast Fourier Transformation: Data Flow

- Basic operation: Butterfly
- Iterated within $\log_2(n)$ stages
- Requires twiddle factors ω^i
 - Roots of unity in \mathbb{C}



Fast Fourier Transformation: Data Flow

- Basic operation: Butterfly
- Iterated within $\log_2(n)$ stages
- Requires twiddle factors ω^i
 - Roots of unity in \mathbb{C}



Fast Fourier Transformation: Complex Arithmetic

- Complex number arithmetic: Requires a floating-point unit!
 - ➡ Challenging to efficiently implement FPU
 - Needs to support complex addition and multiplication
 - Must fit on small area
 - Optimized for FFT computation

Fast Fourier Transformation: Complex Arithmetic

- Complex number arithmetic: Requires a floating-point unit!
 - ➔ Challenging to efficiently implement FPU
 - Needs to support complex addition and multiplication
 - Must fit on small area
 - Optimized for FFT computation

Fast Fourier Transformation: Complex Arithmetic

- Complex number arithmetic: Requires a floating-point unit!
 - ➔ Challenging to efficiently implement FPU
 - Needs to support complex addition and multiplication
 - Must fit on small area
 - Optimized for FFT computation

Fast Fourier Transformation: Complex Arithmetic

- Complex number arithmetic: Requires a floating-point unit!
 - ➔ Challenging to efficiently implement FPU
 - Needs to support complex addition and multiplication
 - Must fit on small area
 - Optimized for FFT computation

Fast Fourier Transformation: Optimizations

- Sharing of resources within the FFT butterfly
 - ➡ lowering area consumption
- Stored *or* on-the-fly generated twiddle factors
 - ➡ increasing flexibility
- Reducing the number of stored twiddle factors by 75%
 - ➡ lowering BRAM consumption

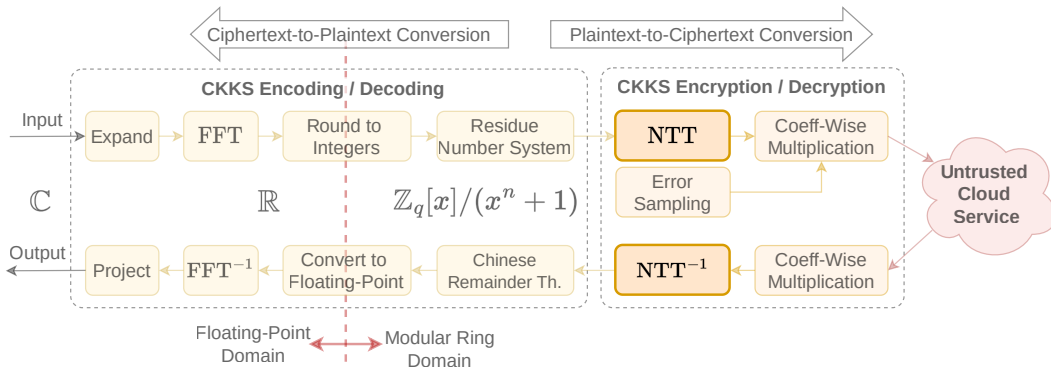
Fast Fourier Transformation: Optimizations

- Sharing of resources within the FFT butterfly
 - ➡ lowering area consumption
- Stored *or* on-the-fly generated twiddle factors
 - ➡ increasing flexibility
- Reducing the number of stored twiddle factors by 75%
 - ➡ lowering BRAM consumption

Fast Fourier Transformation: Optimizations

- Sharing of resources within the FFT butterfly
 - ➡ lowering area consumption
- Stored *or* on-the-fly generated twiddle factors
 - ➡ increasing flexibility
- Reducing the number of stored twiddle factors by 75%
 - ➡ lowering BRAM consumption

Number-Theoretic Transformation



Number-Theoretic Transformation


- Operates over modular ring domain \mathbb{Z}_q
 - ➔ Ring arithmetic modulo prime q
- Special properties of NTT:
 - ➔ Improves runtime of polynomial multiplication:
 $\mathcal{O}(n^2) \rightarrow \mathcal{O}(n \log(n))$
 - ➔ Commonly used in post-quantum cryptography

Number-Theoretic Transformation

- Operates over modular ring domain \mathbb{Z}_q
 - ➡ Ring arithmetic modulo prime q
- Special properties of NTT:
 - ➡ Improves runtime of polynomial multiplication:
 $\mathcal{O}(n^2) \rightarrow \mathcal{O}(n \log(n))$
 - ➡ Commonly used in post-quantum cryptography


Exploiting Similarities of FFT and NTT

- Execution flow of NTT is identical to FFT
 - ➡ Share control logic
 - ➡ Share load and store logic
- Just underlying arithmetic is different
 - ➡ But still allows resource sharing!



Exploiting Similarities of FFT and NTT

- Execution flow of NTT is identical to FFT
 - ➡ Share control logic
 - ➡ Share load and store logic
- Just underlying arithmetic is different
 - ➡ But still allows resource sharing!



Sharing of Integer Multipliers

- FFT instantiates four floating-point multipliers

$$(a_r + ia_i)(b_r + ib_i) = (a_r b_r - a_i b_i) + i(a_r b_i + a_i b_r)$$

➔ each contains one integer multiplier



Use the existing integer multipliers in NTT!

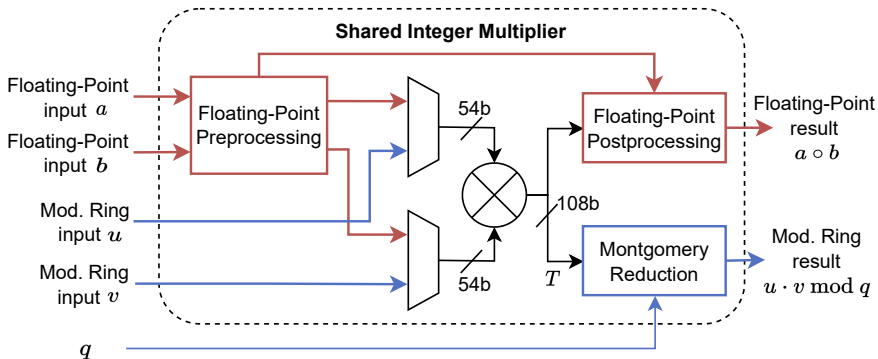
Sharing of Integer Multipliers

- FFT instantiates four floating-point multipliers

$$(a_r + ia_i)(b_r + ib_i) = (a_r b_r - a_i b_i) + i(a_r b_i + a_i b_r)$$

- ➡ each contains one integer multiplier
- 💡 Use the existing integer multipliers in NTT!

Sharing of Integer Multipliers



Sharing of Integer Multipliers: Benefits

- 1 complex multiplier \leftrightarrow 4 modular ring multipliers
- Reused for 3 concurrent NTT + 1 twiddle factor generation
 - ➡ Saves 44% of DSPs

Sharing of Integer Multipliers: Benefits

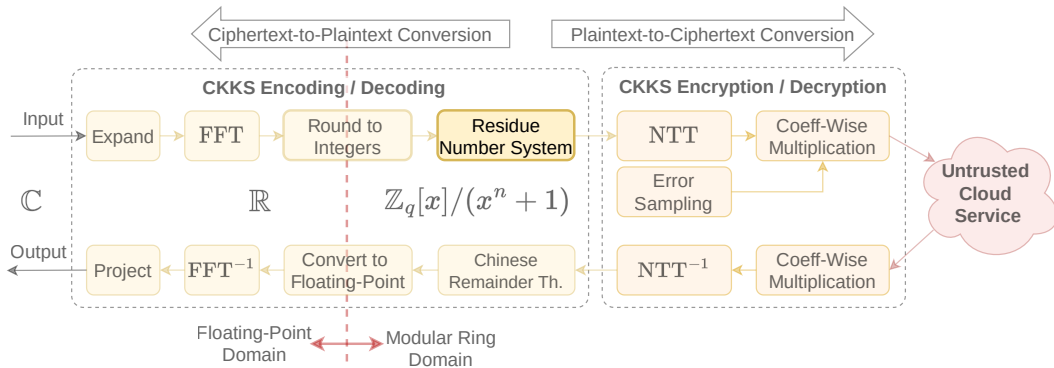
- 1 complex multiplier \leftrightarrow 4 modular ring multipliers
- Reused for 3 concurrent NTT + 1 twiddle factor generation

➡ Saves 44% of DSPs

Sharing of Integer Multipliers: Benefits

- 1 complex multiplier \leftrightarrow 4 modular ring multipliers
- Reused for 3 concurrent NTT + 1 twiddle factor generation
 - ➔ Saves 44% of DSPs

Residue Number System



Residue Number System (in short RNS)

- Reduces computational complexity
- Splits large encoded value x into multiple smaller values (x_0, \dots, x_{L-1})

$$x_i \equiv x \bmod q_i$$

- Challenging to find an efficient hardware solution
 - Large operand sizes x ; small q_i

Residue Number System (in short RNS)

- Reduces computational complexity
- Splits large encoded value x into multiple smaller values (x_0, \dots, x_{L-1})

$$x_i \equiv x \bmod q_i$$

- Challenging to find an efficient hardware solution
 - Large operand sizes x ; small q_i

Residue Number System (in short RNS)

- Reduces computational complexity
- Splits large encoded value x into multiple smaller values (x_0, \dots, x_{L-1})

$$x_i \equiv x \bmod q_i$$

- Challenging to find an efficient hardware solution
 - Large operand sizes x ; small q_i

Residue Number System: Our solution

Input x (202 bits)

mod

Prime q_i (46 bits)

Residue Number System: Our solution


Input x (202 bits)

Comes from Floating-Point: $s \cdot 1.m \cdot 2^e$

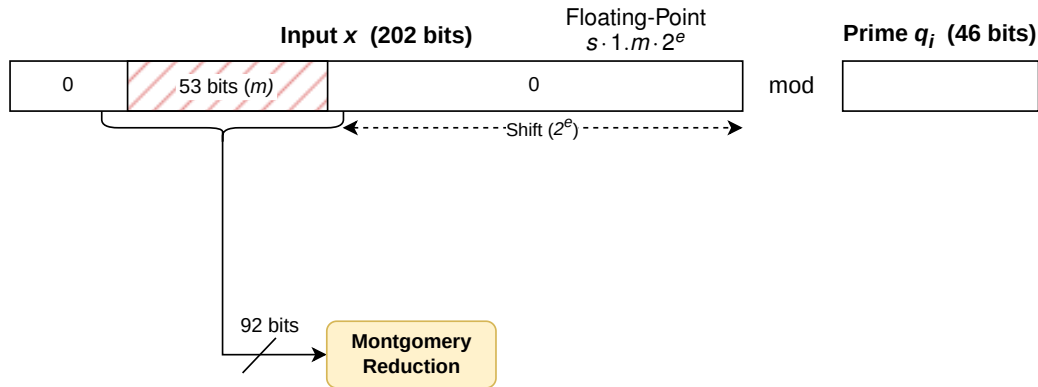
Prime q_i (46 bits)

mod

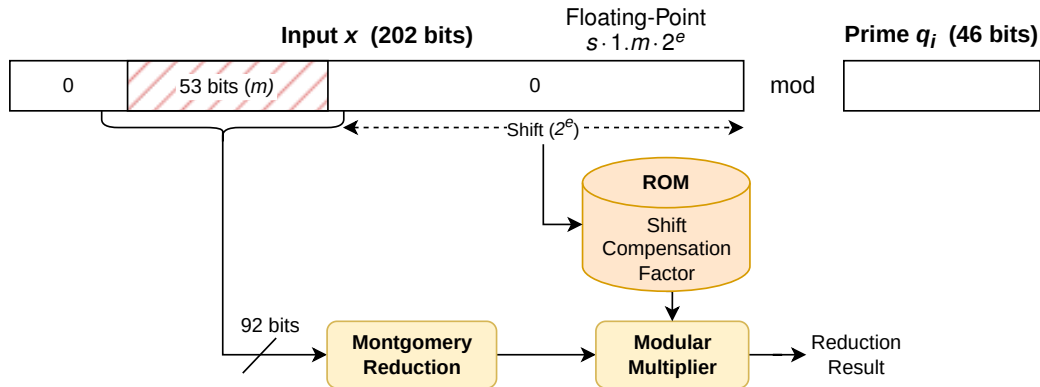
Residue Number System: Our solution



Residue Number System: Our solution



Residue Number System: Our solution



Residue Number System: Our solution

- Benefits compared to existing approach:
 - ➔ Only reuses existing datapath resources
 - ➔ Reduces area consumption significantly
 - ➔ Low latency
 - ➔ Scales well for different parameter sets

Residue Number System: Our solution

- Benefits compared to existing approach:
 - ➔ Only reuses existing datapath resources
 - ➔ Reduces area consumption significantly
 - ➔ Low latency
 - ➔ Scales well for different parameter sets

Residue Number System: Our solution

- Benefits compared to existing approach:
 - ➔ Only reuses existing datapath resources
 - ➔ Reduces area consumption significantly
 - ➔ Low latency
 - ➔ Scales well for different parameter sets

Residue Number System: Our solution

- Benefits compared to existing approach:
 - ➔ Only reuses existing datapath resources
 - ➔ Reduces area consumption significantly
 - ➔ Low latency
 - ➔ Scales well for different parameter sets

Table of Contents

1 Motivation for Homomorphic Encryption

2 Our Proposed Design

3 Implementation Results & Comparison

4 Conclusion

Implementation Results

- Implemented for low-cost FPGAs
 - ZYNQ-7000 FPGA on PYNQ-z2 board
 - 130MHz
 - Kintex-7 FPGA on Genesys2 board
 - 200MHz
- Tested and verified on real hardware

Implementation Results

- Implemented for low-cost FPGAs
 - ZYNQ-7000 FPGA on PYNQ-z2 board
 - 130MHz
 - Kintex-7 FPGA on Genesys2 board
 - 200MHz
- Tested and verified on real hardware


Implementation Results

- Implemented for low-cost FPGAs
 - ZYNQ-7000 FPGA on PYNQ-z2 board
 - 130MHz
 - Kintex-7 FPGA on Genesys2 board
 - 200MHz
- Tested and verified on real hardware


Implementation Results

- Implemented for low-cost FPGAs
 - ZYNQ-7000 FPGA on PYNQ-z2 board
 - 130MHz
 - Kintex-7 FPGA on Genesys2 board
 - 200MHz
- Tested and verified on real hardware


Comparison with Prior Work: Encryption Latency




Comparison with Prior Work: Encryption Latency



Comparison with Prior Work: Encryption Latency




Comparison with Prior Work: Encryption Latency




Performance Figures: Encryption

Our accelerator (FPGA)	Intel Core CPU	Improvement
200 MHz	2.3 GHz	11.5× lower
534 Enc/s	204 Enc/s	2.6×
5.8 mJ/Enc	73.5 mJ/Enc	12.7×



Performance Figures: Encryption

Our accelerator (FPGA)	Intel Core CPU	Improvement
200 MHz	2.3 GHz	11.5× lower
534 Enc/s	204 Enc/s	2.6×
5.8 mJ/Enc	73.5 mJ/Enc	12.7×



Performance Figures: Encryption

Our accelerator (FPGA)	Intel Core CPU	Improvement
200 MHz	2.3 GHz	11.5× lower
534 Enc/s	204 Enc/s	2.6×
5.8 mJ/Enc	73.5 mJ/Enc	12.7×




Table of Contents

1 Motivation for Homomorphic Encryption

2 Our Proposed Design

3 Implementation Results & Comparison

4 Conclusion

Main Contributions

- First hardware accelerator covering all client operations in CKKS
- Unifies the FFT and NTT to save resources
- Novel hardware-friendly approach for RNS
- Achieves performance gain of up to $59\times$
- We made Aloha-HE's source code public on GitHub

Main Contributions

- First hardware accelerator covering all client operations in CKKS
- Unifies the FFT and NTT to save resources
- Novel hardware-friendly approach for RNS
- Achieves performance gain of up to $59\times$
- We made Aloha-HE's source code public on GitHub

Main Contributions

- First hardware accelerator covering all client operations in CKKS
- Unifies the FFT and NTT to save resources
- Novel hardware-friendly approach for RNS
- Achieves performance gain of up to 59×
- We made Aloha-HE's source code public on GitHub

Main Contributions

- First hardware accelerator covering all client operations in CKKS
- Unifies the FFT and NTT to save resources
- Novel hardware-friendly approach for RNS
- Achieves performance gain of up to $59\times$
- We made Aloha-HE's source code public on GitHub

Main Contributions

- First hardware accelerator covering all client operations in CKKS
- Unifies the FFT and NTT to save resources
- Novel hardware-friendly approach for RNS
- Achieves performance gain of up to $59\times$
- We made Aloha-HE's source code public on GitHub

Aloha-HE

A Low-Area Hardware Accelerator for
Client-Side Operations in Homomorphic Encryption

Florian Krieger, Florian Hirner, Ahmet Can Mert, Sujoy Sinha Roy
Institute of Applied Information Processing and Communications, TU Graz

January 31, 2024

References

- CKKS Scheme: J. H. Cheon, A. Kim, M. Kim, and Y. Song, “Homomorphic encryption for arithmetic of approximate numbers,” in *Advances in Cryptology – ASIACRYPT 2017*, T. Takagi and T. Peyrin, Eds. Cham: Springer International Publishing, 2017, pp. 409–437.
- Microsoft SEAL: “Microsoft SEAL (release 4.1),” <https://github.com/Microsoft/SEAL>, Jan 2023, microsoft Research, Redmond, WA.
- SEAL-Embedded: D. Natarajan and W. Dai, “Seal-embedded: A homomorphic encryption library for the internet of things,” *IACR Trans. on Cryptographic Hardware and Embedded Systems*, vol. 2021, no. 3, p. 756–779, Jul. 2021.
- RACE: Z. Azad et al., “Race: Risc-v soc for en/decryption acceleration on the edge for homomorphic computation,” in *Proceedings of the ACM/IEEE International Symposium on Low Power Electronics and Design*. Association for Computing Machinery, 2022.
- RISE: Z. Azad et al., “Rise: Risc-v soc for en/decryption acceleration on the edge for homomorphic encryption,” *IEEE Transactions on Very Large Scale Integration (VLSI) Systems*, pp. 1–14, 2023.
- Matteo *et al.*: S. D. Matteo, M. L. Gerfo, and S. Saponara, “Vlsi design and fpga implementation of an ntt hardware accelerator for homomorphic seal-embedded library,” *IEEE Access*, vol. 11, 2023